

Slurm in 5 min.

Robert Dyro
Stanford ASL

August 15, 2023

Outline

Introduction

Using Slurm

Stanford Sherlock

Introduction - Overview

- ▶ a cluster allows to run independent processes across many computers
- ▶ many computers usually means 2 to 20
- ▶ by all means run your computation on only 1 computer if you want!
- ▶ processes are independent, can communicate over network
- ▶ usage through shell through `ssh`
- ▶ a cluster is managed by a cluster software - Slurm
- ▶ computers in the cluster communicate via network (sockets)
- ▶ storage is shared, all computers see the same files
- ▶ sharing storage is managed by the shared storage software
- ▶ no `sudo`, install locally or load available packages

Introduction - Terminology

- ▶ *a node* - a computer on the network managed by Slurm [-N]
- ▶ *a task* - a single of the many processes you want to run [-n]
- ▶ *a CPU* - a single core on a computer (a physical core) [-c]
- ▶ *a job* - a set of independent processes to run
- ▶ *to submit* - to ask Slurm to run a job
- ▶ *to schedule* - what Slurm does to execute all your tasks
- ▶ *array job* - an alternative way of scheduling NOT covered here

Introduction - How it works

- ▶ you submit a job
 - your job consists of instruction how many and what processes to run
 - you specify strict requirements: nodes nb, memory per node, time limit
 - you specify other requirements and job hints
- ▶ your jobs waits in a queue until resources (nodes) are available
 - if it is rejected, it is rejected immediately
- ▶ your job is scheduled
 - you can get an email about it
 - you check on the status by examining a combined output file
- ▶ your job finishes or is cancelled (by you)

Outline

Introduction

Using Slurm

Stanford Sherlock

Important Commands

Slurm user commands (you're the user, as opposed to an admin)

- ▶ `sinfo` - query available computational resources
- ▶ `squeue` - check status of jobs (`squeue --user $USER`)
- ▶ `srun` - run a single process, low-level command
- ▶ `sbatch` - run a job, a job is a shell script (that calls `srun`)

`module` - allows to load packages, no `sudo apt` on a cluster

- ▶ `module avail` - available preinstalled packages
- ▶ `module list` - currently loaded packages
- ▶ `module load package` - load *package*
- ▶ `module unload package` - unload *package*

Example Job Script

run by issuing `sbatch my_job.sh`

```
#!/bin/bash
#SBATCH --output=out_%j.txt
#SBATCH --time=0-10:00
#SBATCH --mail-type=ALL
#SBATCH --ntasks=20
#SBATCH --cpus-per-task=2
module load julia # load needed packages
source environment # like a python virtualenv
echo ""; date; echo ""
srun -n1 -c2 python3 setup.py # 1 node, task and cpu
wait # barrier #####
for i in {0..100}; do # 101 times, more than ntasks is OK
    srun -n1 -c2 python3 exp.py $i & # notice &
done # each srun ends with & to run all in the background
wait # barrier #####
echo ""; date; echo ""
```


Scheduling

- ▶ when you connect, you're on the login node, don't run stuff there
- ▶ `srun --pty bash` - to run a computational node interactively
- ▶ `srun -p gpu --gpus 1 --pty bash` - a gpu node interactively
- ▶ `#SBATCH -p gpu` in the `my_job.sh` file to have GPU nodes
- ▶ `tail -n 100 out_43598734.txt` to check on a job

Outline

Introduction

Using Slurm

Stanford Sherlock

Sherlock - Stanford University Cluster

All info at: www.sherlock.stanford.edu

- ▶ `ssh $USER@login.sherlock.stanford.edu`
- ▶ <https://login.sherlock.stanford.edu> - web-based interface

Comments

- ▶ uses Slurm
- ▶ access via a linux shell without `sudo` privileges
- ▶ has 100s general computational nodes and ~ 100 GPU nodes
- ▶ account is free, but requires Sherlock admin's and your PI's approval
- ▶ usage is free for research, requires publication acknowledgement
- ▶ can request all the nodes, but will spend time in queue
- ▶ simple job wait: 1 node - 10 s, 10 nodes - 5 min, 20 nodes - 1 hour

Sherlock Storage

Storage

- ▶ three types of storage - *long*, *short* and *very short* term
- ▶ storage is very generous (a lot of GB or TB)
- ▶ *short* and *very short* storage locations are vastly faster
- ▶ *short* storage locations have file expiration (90 days)
- ▶ *very short* storage location is cleared on job end

Available storage locations - these are bash variables

- ▶ *long* term storage
 - `$HOME` - personal long term storage
 - `$GROUP_HOME` - shared research group storage (your lab)
- ▶ *short* term storage
 - `$SCRATCH` - personal short term storage
 - `$GROUP_SCRATCH` - shared research group storage (your lab)
- ▶ *very short* term storage - very good performance
 - `$L_SCRATCH` - local to node
 - `$L_SCRATCH_JOB` - local to node & job

Sherlock Tips

Persistent login

Under `~/.ssh/config` put

```
Host login.sherlock.stanford.edu
  ControlMaster auto
  ControlPath ~/.ssh/%l%r@%h:%p
```

Multiple terminal windows at computational node

- ▶ query what nodes are yours `squeue --user $USER`
- ▶ issue `ssh sh02-01n43`

Installing things

- ▶ Sherlock does not give `sudo` access
- ▶ many installations do not *really* require `sudo`
- ▶ install and use stuff from `~/.local/lib` `~/.local/bin`

Copying files back and forth

- ▶ use `rsync`! excellent tutorial from Digital Ocean [here](#)